

Thema09_EDA

2023-01-29

Introduction

Pancreatic ductal adenocarcinoma (PDAC) is a severe type of cancer with a poor prognosis. The current five-year survival rate is less than 10%. The main reason for this poor outcome is the lack of effective methods for early detection. In most cases, PDAC is not diagnosed until it has reached an advanced stage, at which point treatment options are limited and the chances of survival are greatly reduced.

The early detection of PDAC is crucial for improving patient outcomes, but current diagnostic methods are inadequate. Imaging techniques such as CT and MRI are not reliable for early detection and biomarkers for PDAC have been hard to identify. However, recent research suggests that urinary biomarkers may be useful for early detection of PDAC.

This report presents a machine learning approach for the early detection of PDAC using urinary biomarkers such as creatinine, LYVE1, REG1A, REG1B, and TFF1. A Java Wrapper is used to predict if a patient has PDAC, benign pancreatic conditions, or no PDAC. The data for this study was obtained from Kaggle and analyzed using RStudio Markdown. The goal of this report is to investigate the research question: “Is it possible to detect pancreatic cancer using values of the urinary biomarkers?” and demonstrate the potential usefulness of this machine learning approach for early detection of PDAC and to provide insight into the biomarkers that can be used for this purpose. By providing a way to detect PDAC early, this machine learning approach could help improve patient outcomes and increase the chances of survival.

EDA

Codebook

The first step of the EDA was to create a codebook of the data to visualize the data and understand the attributes of the dataset. This codebook includes the column names, data types, units of measurement, and a brief description of each column. This helped in understanding the data and identifying the relevant columns for the analysis.

```
myData <- read.csv("Data/Debernardi et al 2020 data.csv")

columns <- colnames(myData)
type <- c("character", "character", "character", "double", "character", "double", "logical", "logical",
unit <- c(NA, NA, NA, "years", "F/M", NA, NA, NA, "U/ml", "mg/ml", "ng/ml", "ng/ml", "ng/ml", "ng/ml")
descriptions = c("Unique string identifying each subject", "Cohort 1, previously used samples; Cohort 2
codebook <- data.frame(columns, type, unit, descriptions)
write.csv(codebook, "Codebook.csv", row.names = FALSE)
```

```
pander(codebook, booktabs = T, caption="The Codebook")
```

Table 1: The Codebook

columns	type	unit	descriptions
sample_id	character	NA	Unique string identifying each subject
patient_cohort	character	NA	Cohort 1, previously used samples; Cohort 2, newly added samples
sample_origin	character	NA	BPTB: Barts Pancreas Tissue Bank, London, UK; ESP: Spanish National Cancer Research Centre, Madrid, Spain; LIV: Liverpool University, UK; UCL: University College London, UK
age	double	years	Age in years
sex	character	F/M	M = male, F = female
diagnosis	double	NA	1 = control (no pancreatic disease), 2 = benign hepatobiliary disease (119 of which are chronic pancreatitis); 3 = Pancreatic ductal adenocarcinoma, i.e. pancreatic cancer
stage	logical	NA	For those with pancreatic cancer, what stage was it? One of IA, IB, IIA, IIIB, III, IV
benign_sample_diagnosis	logical	NA	For those with a benign, non-cancerous diagnosis, what was the diagnosis?
plasma_CA19_9	double	U/ml	Blood plasma levels of CA 19-9 monoclonal antibody that is often elevated in patients with pancreatic cancer. Only assessed in 350 patients (one goal of the study was to compare various CA 19-9 cutpoints from a blood sample to the model developed using urinary samples).
creatinine	double	mg/ml	Urinary biomarker of kidney function
LYVE1	double	ng/ml	Urinary levels of Lymphatic vessel endothelial hyaluronan receptor 1, a protein that may play a role in tumor metastasis
REG1B	double	ng/ml	Urinary levels of a protein that may be associated with pancreas regeneration.
TFF1	double	ng/ml	Urinary levels of Trefoil Factor 1, which may be related to regeneration and repair of the urinary tract
REG1A	double	ng/ml	Urinary levels of a protein that may be associated with pancreas regeneration. Only assessed in 306 patients (one goal of the study was to assess REG1B vs REG1A)

Data exploration

Visualization

The table created in the EDA process provided a clear overview of the distribution of the data for the relevant columns. The diagnosis column showed the diagnosis of the sample, with 1 indicating no PDAC, 2 indicating benign hepatobiliary disease (non-cancerous, non-harmful pancreatic condition), and 3 indicating that the sample has PDAC. By viewing the distribution of this column, it was possible to identify the proportion of cases with PDAC and benign pancreatic conditions in the dataset.

The age column provided information about the age of the patient at the time of sample collection. This information can be used to identify if there is a correlation between age and the development of PDAC. The sex column provided information about the gender of the patient, which can also be used to identify if there is a correlation between gender and the development of PDAC.

The creatinine, LYVE1, REG1A, REG1B, and TFF1 columns provided information about the levels of the biomarkers in the urine samples. By viewing the distribution of these columns, it was possible to identify the range of values for each biomarker and if there is a difference between the levels of the biomarkers in cases of PDAC and benign pancreatic conditions. The table also helped to identify if there are any missing values in the data, which may have to be handled before applying the machine learning algorithms.

Overall, the table provided a comprehensive overview of the data, allowing for a better understanding of the dataset and identifying any potential issues that may need to be addressed before applying the machine learning algorithms. It also helped in identifying if there are any correlation between the variables which can help us to select the best quality metric and machine learning algorithm.

```
relevant <- myData[c(4, 6, 9:14)]  
pander(head(relevant, n=15), booktabs = T, caption = "Values of biomarkers and the diagnosis")
```

Table 2: Values of biomarkers and the diagnosis

age	diagnosis	plasma_CA19_9	creatinine	LYVE1	REG1B	TFF1	REG1A
33	1	11.7	1.832	0.8932	52.95	654.3	1262
81	1	NA	0.9727	2.038	94.47	209.5	228.4
51	1	7	0.7804	0.1456	102.4	461.1	NA
61	1	8	0.7012	0.002805	60.58	142.9	NA
62	1	9	0.2149	0.0008596	65.54	41.09	NA
53	1	NA	0.8482	0.003393	62.13	59.79	NA
70	1	NA	0.622	0.1744	152.3	117.5	NA
58	1	11	0.8935	0.003574	3.73	40.29	NA
59	1	NA	0.4863	0.001945	7.021	26.78	NA
56	1	24	0.6107	0.2788	83.93	19.18	NA
77	1	NA	0.2941	0.001176	6.218	28.3	NA
71	1	23	1.052	0.8603	243.1	608.3	NA
49	1	NA	0.8596	1.416	151.8	74.19	505.6
53	1	7	1.911	1.517	150.9	590.7	NA
56	1	12	0.9161	0.5996	93.81	93.58	NA

As is it obvious to see, the REG1A column contains a lot of missing values. The question is exactly how is this possible and what does it mean for the rest of the data? The first and most simple solution is that we don't need this column at all, since the REG1B is similar in function and containing all of the data. Another solution that's not very logical, is to use only the rows where there is a value in the REG1A column. The most logical solution in my opinion is to look at every column separately, and then remove the missing data.

When the cleaned data then is plotted, it's possible to look at all the columns together and see if there is any correlation or trend to figure out.

To identify any potential outliers, a boxplot was created using the relevant columns. The boxplot was created using ggplot2, and it helps in identifying any outliers in the data. The outliers were then analyzed and checked if they could be removed or not.

```
p1 <- ggplot(myData, aes(x=age)) +
  geom_histogram(fill = "lightgrey", col = "black", binwidth = 1)

p2 <- ggplot(myData, aes(x=creatinine)) +
  geom_boxplot()
p3 <- ggplot(myData, aes(x=LYVE1)) +
  geom_boxplot()
p4 <- ggplot(myData, aes(x=REG1B)) +
  geom_boxplot()
p5 <- ggplot(myData, aes(x=TFF1)) +
  geom_boxplot()
p6 <- ggplot(myData, aes(x=REG1A)) +
  geom_boxplot()

grid.arrange(p1, p2, p3, p4, p5, p6, nrow = 3, bottom="Figure 1: Biomarkers distribution raw data")
```




Figure 1: Biomarkers distribution raw data

as you can see, there are a lot of outliers which we need to check. We can normalize the data using a log transformation.

```
log <- log(myData[9:14] + 1)
myData[9:14] <- log
```

Before proceeding, it is quite helpful to look at the distribution of the diagnosis column, to see if they're evenly distributed or not.

```
ggplot(myData, aes(x=diagnosis)) +
  geom_histogram(fill = "lightgrey", col = "black", binwidth = 1) +
  ggtitle("Distribution of the diagnosis") +
  labs(caption = "Figure 2: Diagnosis distribution") +
  theme(plot.caption = element_text(size = 14))
```




Figure 2: Diagnosis distribution

What we should do now, is look at the values of those urinary levels with each diagnosis and see if there are any obvious differences.

Let's start with the samples that do not have PDAC

```
myData1 <- subset(myData, myData$diagnosis == 1)

p1 <- ggplot(myData1, aes(x=age)) +
  geom_histogram(fill = "lightgrey", col = "black", binwidth = 1) +
  ggtitle("Age distribution")

p2 <- ggplot(myData1, aes(x=creatinine)) +
  geom_boxplot() +
  ggtitle("Distribution of creatinine")
p3 <- ggplot(myData1, aes(x=LYVE1)) +
  geom_boxplot() +
  ggtitle("Distribution of LYVE1")
p4 <- ggplot(myData1, aes(x=REG1B)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1B")
p5 <- ggplot(myData1, aes(x=TFF1)) +
  geom_boxplot() +
  ggtitle("Distribution of TFF1")
p6 <- ggplot(myData1, aes(x=REG1A)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1A")

grid.arrange(p1, p2, p3, p4, p5, p6, nrow = 3, bottom="Figure 3: Biomarkers distribution of samples with")
```




Figure 3: Biomarkers distribution of samples with PDAC

Here are the values of the samples with non-cancerous pancreatic conditions.

```
myData2 <- subset(myData, myData$diagnosis == 2)

p1 <- ggplot(myData2, aes(x=age)) +
  geom_histogram(fill = "lightgrey", col = "black", binwidth = 1) +
  ggtitle("Age distribution")

p2 <- ggplot(myData2, aes(x=creatinine)) +
  geom_boxplot() +
  ggtitle("Distribution of creatinine")
p3 <- ggplot(myData2, aes(x=LYVE1)) +
  geom_boxplot() +
  ggtitle("Distribution of LYVE1")
p4 <- ggplot(myData2, aes(x=REG1B)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1B")
p5 <- ggplot(myData2, aes(x=TFF1)) +
  geom_boxplot() +
  ggtitle("Distribution of TFF1")
p6 <- ggplot(myData2, aes(x=REG1A)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1A")

grid.arrange(p1, p2, p3, p4, p5, p6, nrow = 3, bottom="Figure 4: Biomarkers distribution of samples with
```




Figure 4: Biomarkers distribution of samples with benign pancreatic conditions

And finally the samples that do have PDAC.

```
myData3 <- subset(myData, myData$diagnosis == 3)

p1 <- ggplot(myData3, aes(x=age)) +
  geom_histogram(fill = "lightgrey", col = "black", binwidth = 1) +
  ggtitle("Age distribution")

p2 <- ggplot(myData3, aes(x=creatinine)) +
  geom_boxplot() +
  ggtitle("Distribution of creatinine")
p3 <- ggplot(myData3, aes(x=LYVE1)) +
  geom_boxplot() +
  ggtitle("Distribution of LYVE1")
p4 <- ggplot(myData3, aes(x=REG1B)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1B")
p5 <- ggplot(myData3, aes(x=TFF1)) +
  geom_boxplot() +
  ggtitle("Distribution of TFF1")
p6 <- ggplot(myData3, aes(x=REG1A)) +
  geom_boxplot() +
  ggtitle("Distribution of REG1A")

grid.arrange(p1, p2, p3, p4, p5, p6, nrow = 3, bottom="Figure 5: Biomarkers distribution of samples with PDAC")
```

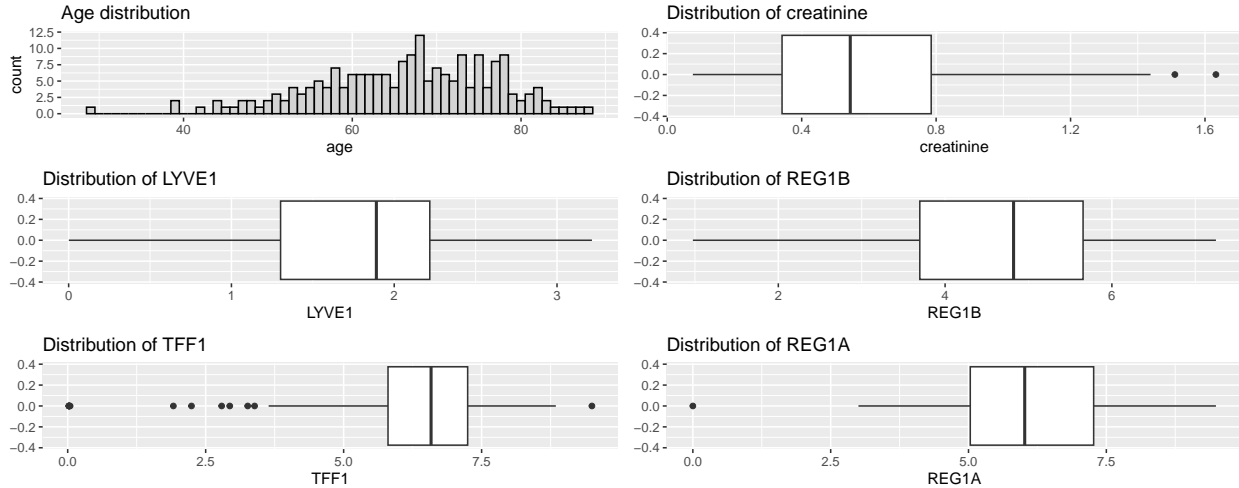


Figure 5: Biomarkers distribution of samples without PDAC

At first glance, there are notable differences but we can further investigate this in future steps of this research.

```
write.csv(myData, "Data/DataCleaned.csv")
```

Correlation

The correlation data obtained during the EDA process helped to investigate if the biomarkers correlated with each other in any way. By understanding the correlation between the biomarkers, it was possible to identify which biomarkers could be used as predictors of others.

By analyzing the correlation matrix, it was possible to identify if any of the biomarkers were highly correlated with others. For example, if creatinine and LYVE1 had a high positive correlation coefficient, it would indicate that when the creatinine levels were high, the LYVE1 levels were also high. This information can be used to identify which biomarkers could be used as predictors of others and which biomarkers are more informative.

Additionally, it also helped to identify if there is any multicollinearity in the data. Multicollinearity occurs when two or more independent variables in a regression analysis are highly correlated. This can lead to unstable and unreliable estimates of the regression coefficients and can also cause problems in interpreting the results of the analysis. By identifying any multicollinearity, it was possible to address it and ensure that the machine learning algorithm was not affected by it.

Overall, the correlation data helped in identifying which biomarkers are more informative, which biomarkers could be used as predictors of others, and if there is any multicollinearity in the data. This information was used to select the best quality metric and machine learning algorithm.

```
p1 <- ggplot(myData, aes(x=creatinine, y=LYVE1)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
  ggtitle("Creatinine vs LYVE1")
p2 <- ggplot(myData, aes(x=creatinine, y=REG1B)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
  ggtitle("Creatinine vs REG1B")
p3 <- ggplot(myData, aes(x=creatinine, y=TFF1)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
```



```

ggtitle("Creatinine vs TFF1")
p4 <- ggplot(myData, aes(x=LYVE1, y=REG1B)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
  ggtitle("LYVE1 vs REG1B")
p5 <- ggplot(myData, aes(x=LYVE1, y=TFF1)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
  ggtitle("LYVE1 vs TFF1")
p6 <- ggplot(myData, aes(x=TFF1, y=REG1B)) +
  geom_point() +
  geom_smooth(method='lm', formula = y~x) +
  ggtitle("TFF1 vs REG1B")

```

grid.arrange(p1, p2, p3, p4, p5, p6, nrow=3, bottom="Figure 6: Scatter plot showing the correlation between

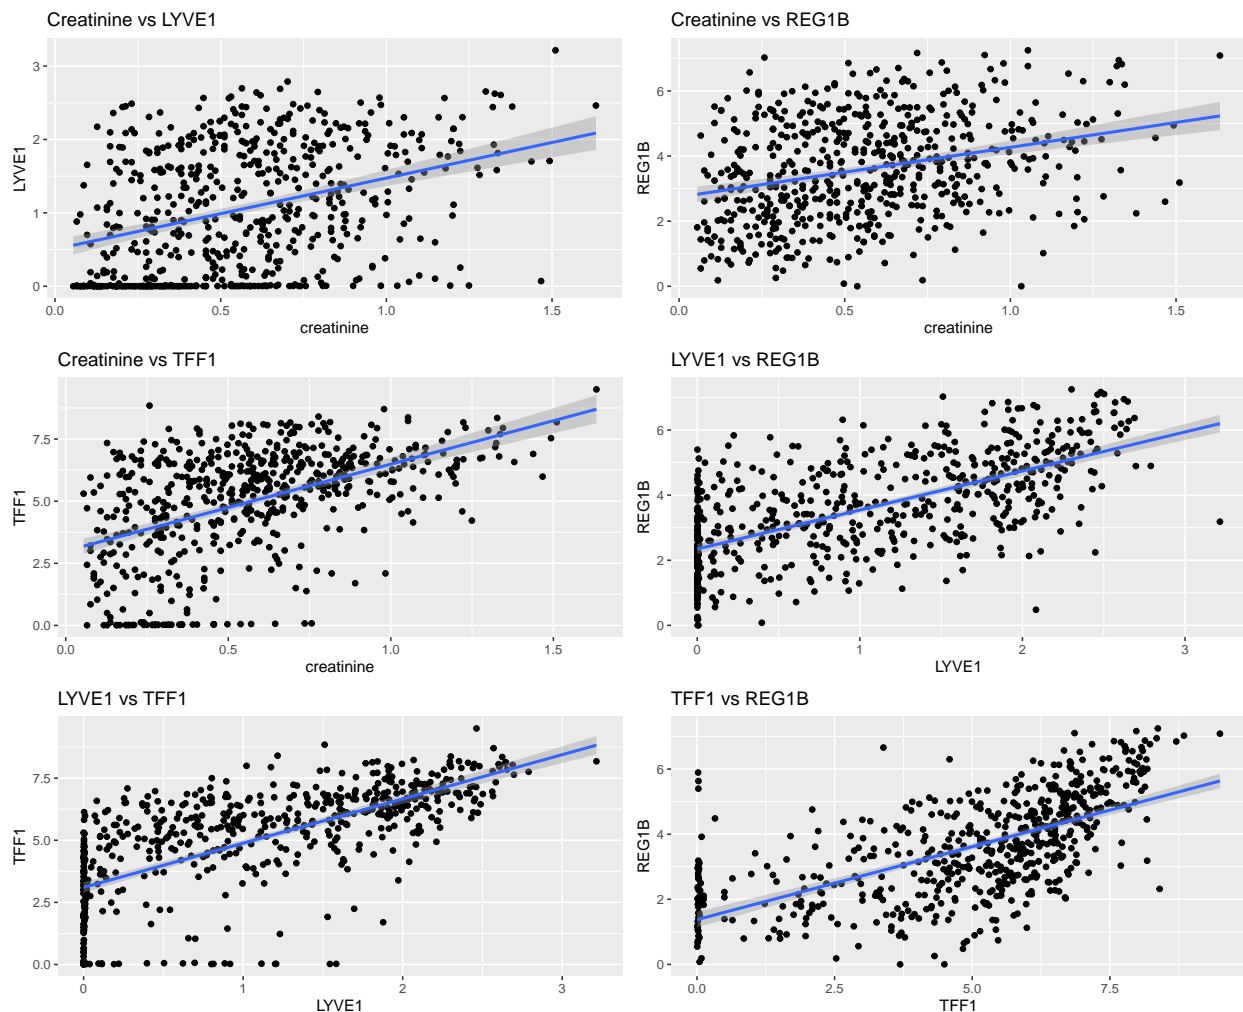


Figure 6: Scatter plot showing the correlation between creatinine, LYVE1, REG1A, REG1B, and TFF1 biomarkers in urinary samples of pancreatic cancer patients

In this visualization there is a bit of correlation visible

Now, let's prepare the data for machine learning. The first step will be to remove the columns that aren't really necessary and relevant to the research to make it more clear.

```
myData <- subset(myData, select = -c(patient_cohort, sample_origin, benign_sample_diagnosis))
pander(summary(myData), booktabs = T, caption = "Relevant data")
```

Table 3: Relevant data (continued below)

sample_id	age	sex	diagnosis
Length:590	Min. :26.00	Length:590	Min. :1.000
Class :character	1st Qu.:50.00	Class :character	1st Qu.:1.000
Mode :character	Median :60.00	Mode :character	Median :2.000
NA	Mean :59.08	NA	Mean :2.027
NA	3rd Qu.:69.00	NA	3rd Qu.:3.000
NA	Max. :89.00	NA	Max. :3.000
NA	NA	NA	NA

Table 4: Table continues below

stage	plasma_CA19_9	creatinine	LYVE1
Length:590	Min. : 0.000	Min. :0.05501	Min. :0.000129
Class :character	1st Qu.: 2.197	1st Qu.:0.31717	1st Qu.:0.154584
Mode :character	Median : 3.314	Median :0.54455	Median :0.974503
NA	Mean : 3.882	Mean :0.56738	Mean :1.054778
NA	3rd Qu.: 5.687	3rd Qu.:0.76056	3rd Qu.:1.825361
NA	Max. :10.342	Max. :1.63254	Max. :3.214479
NA	NA's :240	NA	NA

REG1B	TFF1	REG1A
Min. :0.001104	Min. :0.005279	Min. :0.000
1st Qu.:2.464467	1st Qu.:3.805664	1st Qu.:4.403
Median :3.563973	Median :5.563994	Median :5.345
Mean :3.607996	Mean :4.979594	Mean :5.411
3rd Qu.:4.818184	3rd Qu.:6.611675	3rd Qu.:6.477
Max. :7.247720	Max. :9.498920	Max. :9.488
NA	NA	NA's :284